

# CombinedModel

December 6, 2021

## 0.0.1 Classification Models

```
[1]: !pip install -U imbalanced-learn
!pip install xgboost
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import xgboost as xgb
from numpy import std
import time
import statistics
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, \
    mean_squared_error, mean_absolute_error, confusion_matrix, \
    classification_report
from sklearn.model_selection import cross_val_score, KFold, train_test_split
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.compose import make_column_transformer
from sklearn.pipeline import make_pipeline
from collections import Counter
from imblearn.over_sampling import SMOTENC
from scipy.stats import shapiro
```

```
Requirement already up-to-date: imbalanced-learn in
/opt/conda/lib/python3.7/site-packages (0.8.1)
Requirement already satisfied, skipping upgrade: scipy>=0.19.1 in
/opt/conda/lib/python3.7/site-packages (from imbalanced-learn) (1.4.1)
Requirement already satisfied, skipping upgrade: joblib>=0.11 in
/opt/conda/lib/python3.7/site-packages (from imbalanced-learn) (0.15.1)
Requirement already satisfied, skipping upgrade: scikit-learn>=0.24 in
/opt/conda/lib/python3.7/site-packages (from imbalanced-learn) (1.0.1)
Requirement already satisfied, skipping upgrade: numpy>=1.13.3 in
/opt/conda/lib/python3.7/site-packages (from imbalanced-learn) (1.18.4)
Requirement already satisfied, skipping upgrade: threadpoolctl>=2.0.0 in
/opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.24->imbalanced-
learn) (3.0.0)
```

Requirement already satisfied: xgboost in /opt/conda/lib/python3.7/site-packages (1.5.1)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.4.1)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.18.4)

### Import Dataset

```
[2]: import pandas as pd
df = pd.read_csv('Final_Dataset.csv', index_col = 0)
```

### Count value of Stars attribute to verify imbalance

```
[3]: df['stars'].value_counts()
```

```
[3]: 4    5835
     3    1186
     5     302
     2      81
     1      17
     Name: stars, dtype: int64
```

### Classification Algorithms with Imbalanced Dataset

#### Seperate the data into Train and Test

```
[4]: x = df.drop('stars', axis = 1)
     y = df['stars']

     x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.30,
     ↪random_state = 42)
     print(Counter(y_train))
     print(Counter(y_test))
```

```
Counter({4: 4072, 3: 824, 5: 227, 2: 59, 1: 12})
```

```
Counter({4: 1763, 3: 362, 5: 75, 2: 22, 1: 5})
```

```
[5]: def convert(data):
     encode = LabelEncoder()
     data['business_id'] = encode.fit_transform(data.business_id)
     data['city'] = encode.fit_transform(data.city)
     data['state'] = encode.fit_transform(data.state)
     data['categories'] = encode.fit_transform(data.categories)
     data['review_id'] = encode.fit_transform(data.review_id)
     data['date'] = encode.fit_transform(data.date)
```

```

data['text'] = encode.fit_transform(data.text)
data=data.fillna(-999)
return data

x_train = convert(x_train)
x_test = convert(x_test)

```

```
[6]: print(x_train.dtypes)
```

```

business_id    int64
city           int64
state          int64
review_count   int64
is_open        int64
...
veri tasti     int64
veri tasty     int64
wait time      int64
we order       int64
we tri         int64
Length: 185, dtype: object

```

```
[7]: # Training with imbalance data
classifier = LogisticRegression(solver='lbfgs', max_iter=1000)
classifier.fit(x_train, y_train)
print(classification_report(y_test, classifier.predict(x_test)))

hist1 = plt.hist(y_train)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Distribution of Class "Stars" before applying oversample')
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5
2	0.00	0.00	0.00	22
3	0.00	0.00	0.00	362
4	0.79	1.00	0.88	1763
5	0.00	0.00	0.00	75
accuracy			0.79	2227
macro avg	0.16	0.20	0.18	2227
weighted avg	0.63	0.79	0.70	2227

```

/opt/conda/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

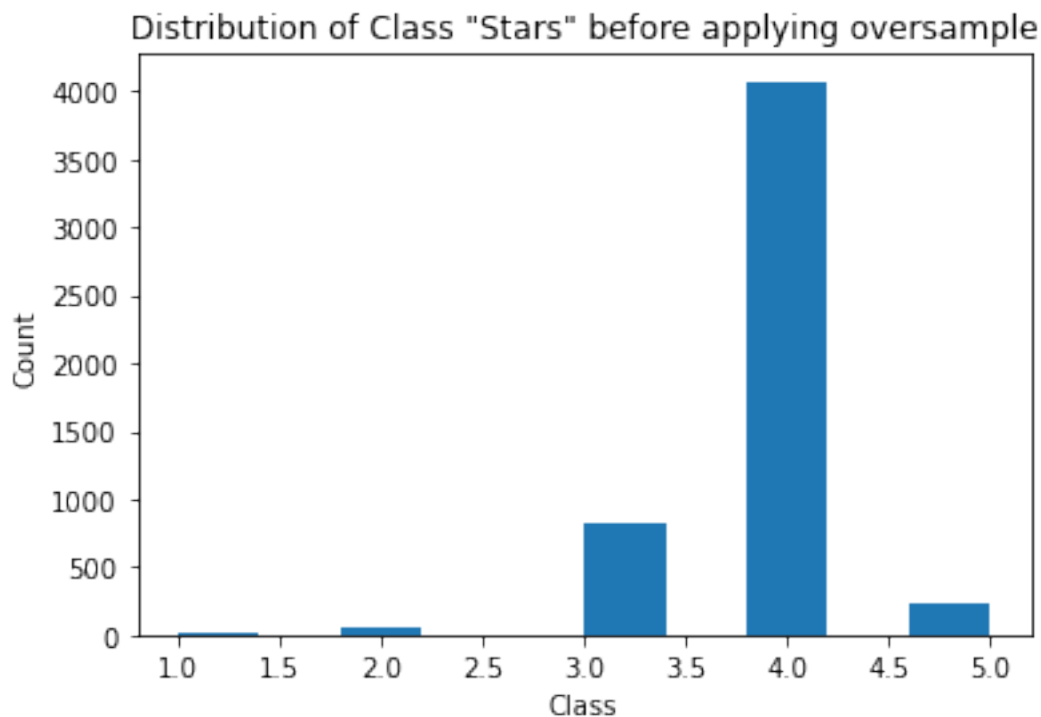
```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))  
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))  
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

[7]: Text(0.5, 1.0, 'Distribution of Class "Stars" before applying oversample')



```
[8]: # Create an oversampled training data
smote = SMOTENC([4,5,*range(10,185,1)], random_state = 42)
x_oversample, y_oversample = smote.fit_resample(x_train, y_train)

# Training with oversampled data
classifier_o = LogisticRegression(solver='lbfgs', max_iter=1000)
classifier_o.fit(x_oversample, y_oversample)
print(classification_report(y_test, classifier_o.predict(x_test)))


hist2 = plt.hist(y_oversample)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Distribution of Class "Stars" after applying oversample')
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5
2	0.05	0.23	0.09	22
3	0.24	0.10	0.14	362
4	0.91	0.56	0.70	1763
5	0.07	0.85	0.13	75
accuracy			0.49	2227
macro avg	0.26	0.35	0.21	2227
weighted avg	0.76	0.49	0.58	2227

/opt/conda/lib/python3.7/site-packages/sklearn/linear\_model/\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
[8]: Text(0.5, 1.0, 'Distribution of Class "Stars" after applying oversample')
```



```
[9]: forest = RandomForestClassifier()
forest.fit(x_train, y_train)
y_pred = forest.predict(x_test)
```

```
[10]: accuracyforest = (accuracy_score(y_test, y_pred)*100)
print('Accuracy is ', str(round(accuracyforest, 3)))
```

Accuracy is 80.242

```
[11]: rooterror = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Square Error(root-mean-square-error):%f" % (rooterror))
```

Root Mean Square Error(root-mean-square-error):0.493787


```
[12]: #View confusion Matrix for test data and predictions
matrix = confusion_matrix(y_test, y_pred)
print(matrix)
```

```
[[ 0  0  0  5  0]
 [ 0  0  4 18  0]
 [ 0  0 41 321 0]
 [ 0  0 18 1745 0]
 [ 0  0  3  71  1]]
```

```
[13]: matrix = confusion_matrix(y_test, y_pred)
matrix = matrix.astype('float')/ matrix.sum(axis = 1)[:, np.newaxis]

plt.figure(figsize = (16, 7))
sns.set(font_scale = 1.4)
sns.heatmap(matrix, annot = True, annot_kws = {'size': 10},
            cmap = plt.cm.Blues, linewidths = 0.2)

class_names = ['1', '2', '3', '4', '5']
tick_marks = np.arange(len(class_names))
tick_marks2 = tick_marks + 0.5
plt.xticks(tick_marks, class_names, rotation=25)
plt.yticks(tick_marks2, class_names, rotation=0)
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix for Random Forest Model')
plt.show()
```



```
[14]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5
2	0.00	0.00	0.00	22
3	0.62	0.11	0.19	362
4	0.81	0.99	0.89	1763
5	1.00	0.01	0.03	75

accuracy			0.80	2227
macro avg	0.49	0.22	0.22	2227
weighted avg	0.77	0.80	0.74	2227

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

### K-Fold Cross Validation for Random Forest

```
[15]: model = RandomForestClassifier(random_state=42)
kf = KFold(n_splits = 5, random_state = 42, shuffle = True)
kf.get_n_splits(x)
i = 1
print("Confusion Matrix:")
for train_index, test_index in kf.split(x):
    x_train, x_test = x.iloc[train_index], x.iloc[test_index]
    y_train, y_test_kfold = y.iloc[train_index], y.iloc[test_index]

    x_train = convert(x_train)
    x_test = convert(x_test)

    x_train_oversampled, y_train_oversampled = smote.fit_resample(x_train,
→y_train)
    model.fit(x_train, y_train)
    print(i, "\n", confusion_matrix(y_test_kfold, model.predict(x_test)))
    i = i + 1
    print(10 * '#')
```

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:3:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-
```



docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
This is separate from the ipykernel package so we can avoid doing imports until

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

after removing the cwd from sys.path.

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

"""

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:6:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:7:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

import sys

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:8:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:9:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
if __name__ == '__main__':
```

Confusion Matrix:

1

```
[[ 0  0  3  1  0]
 [ 0  0  5 10  0]
 [ 0  0 32 215  1]
 [ 0  0 29 1140  0]
 [ 0  0  4  45  0]]
```

#####

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

after removing the cwd from sys.path.

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

"""

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:6:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:7:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
import sys
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:8:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:9:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
if __name__ == '__main__':
2
[[ 0  0  1  0  0]
 [ 0  0  0 17  0]
 [ 0  0 15 199  0]
 [ 0  0 10 1185  0]
 [ 0  0  1  55  1]]
#####
```

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:3:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
after removing the cwd from sys.path.
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:5:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
"""
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:6:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:7:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
import sys
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:8:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:9:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
if __name__ == '__main__':
```

```
3
[[ 0  0  2  5  0]
 [ 0  0  4 12  0]
 [ 0  0 28 204  0]
 [ 0  0 15 1152  0]
 [ 0  0  1  61  0]]
```

```
#####
```

```
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:3:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:4:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

after removing the cwd from sys.path.

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:5:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

"""

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:6:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:7:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

import sys

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:9:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
if __name__ == '__main__':
```

4

```
[[ 0  0  1  2  0]
 [ 0  0  5 15  0]
 [ 0  0 25 249 0]
 [ 0  0 12 1101 0]
 [ 0  0  1  73  0]]
```

#####

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

after removing the cwd from sys.path.

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

"""

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:6:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:7:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
import sys
/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:8:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:9:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
if __name__ == '__main__':
5
[[ 0  0  2  0  0]
 [ 0  0  4  9  0]
 [ 0  0 22 196  0]
 [ 0  0 14 1177  0]
 [ 0  0  3  56  1]]
#####

```

```
[16]: y_pred_kfold = model.predict(x_test)
```

```
[17]: print(classification_report(y_test_kfold, y_pred_kfold))
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	2
2	0.00	0.00	0.00	13
3	0.49	0.10	0.17	218
4	0.82	0.99	0.90	1191
5	1.00	0.02	0.03	60
accuracy			0.81	1484
macro avg	0.46	0.22	0.22	1484
weighted avg	0.77	0.81	0.74	1484

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1308:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[18]: x = convert(x)
score = cross_val_score(model, x, y, cv = kf)
# score
print("%0.2f accuracy with a standard deviation of %0.2f" % (score.mean(),
↪score.std()))
```

0.83 accuracy with a standard deviation of 0.02

```
[ ]:
```